

Hough Transform

Dr Wenjia Bai

Department of Computing & Brain Sciences

Hough transform

- After edge detection, we get a binary edge map, which contains a lot of edge points.
- If these edge points belong to a line, how can we get a parametric representation of the line?

Line parameterisation

- For example, a line can be represented by $y = mx + b$, i.e. just two parameters m and b .
- This is a much more efficient representation than a lot of edge points.

Line parameterisation

- Slope intercept form

$$y = mx + b$$

- Double intercept form

$$\frac{x}{a} + \frac{y}{b} = 1$$


- Normal form

$$x\cos(\theta) + y\sin(\theta) = \rho$$

Slope intercept form

$$y = mx + b$$


- m : slope
- b : y -intercept



Double intercept form

$$\frac{x}{a} + \frac{y}{b} = 1$$

- a : x -intercept
- b : y -intercept



Normal form

$$x\cos(\theta) + y\sin(\theta) = \rho$$


- θ : angle
- ρ : distance from origin

- Derivation

x-intercept $a = \frac{\rho}{\cos(\theta)}$, y-intercept $b = \frac{\rho}{\sin(\theta)}$

Plug into $\frac{x}{a} + \frac{y}{b} = 1$

We can get the normal form.



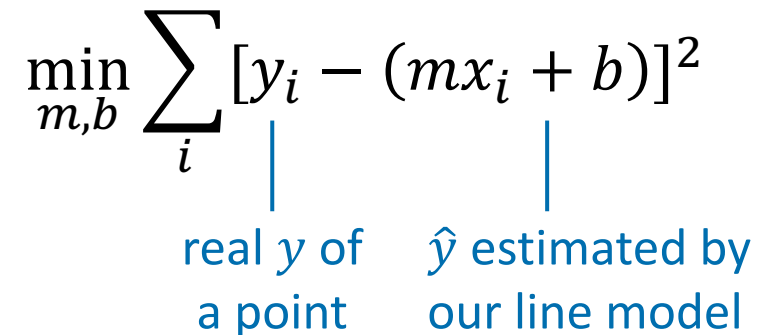
Hough transform

- Hough transform is a transform from image space to parameter space (e.g. from an edge map to the two parameters of a line).
- Its output is a parametric model, given the input edge points.
- The basic idea is that each edge point votes for possible models in the parameter space.

Model fitting

- Some of you may have a different idea here, especially if you know optimisation and model fitting.
- One way to solve this problem is to fit a line model onto the edge points.
 - Suppose we have a set of points $(x_1, y_1), (x_2, y_2), \dots$ and we would like to fit a line model $y = mx + b$ to these points.
 - (m, b) can be estimated by minimising the fitting error

$$\min_{m,b} \sum_i [y_i - (mx_i + b)]^2$$


real y of a point \hat{y} estimated by our line model

- How will Hough transform solve the problem differently?

Hough transform

- Let us use the slope intercept form for a line model

$$y = mx + b$$



$$b = y - mx$$

- We have edge points in the image space $(x_1, y_1), (x_2, y_2), (x_3, y_3) \dots$
- Each point votes for a line model in the parameter space.
- For example, the first point will vote for $b = y_1 - mx_1$.

Hough transform





Image space

$$b = y - mx$$




Parameter space

Vote result:
 $m = 1, b = 0$
 $y = x$

Hough transform

- In practice, the parameter space is divided into 2D bins.
- Each point increments the vote by 1 in one of the bins.
- One problem with the slope intercept form:
 - The parameter space is too large.
 - $m \in [-\infty, +\infty]$, $b \in [-\infty, +\infty]$
 - We need a lot of bins.



Hough transform

- Solution

- Use the normal form instead

$$x\cos(\theta) + y\sin(\theta) = \rho$$

- Although $\rho \in [-\infty, +\infty]$, at least $\theta \in [0, \pi)$.
 - We can use much fewer bins.
 - By the way, in practice ρ is not infinite either. It is limited by the image size.
- The transform from image space to parameter space will look different.

Hough transform





Image space

$$x\cos(\theta) + y\sin(\theta) = \rho$$



θ goes from 0 to π .

Parameter space

Vote result:
 $\theta = \frac{3}{4}\pi, \rho = 0$

Line detection by Hough transform

Algorithm

Initialise the bins $H(\rho, \theta)$ to all zeros.

For each edge point (x, y)

 For θ from 0 to π

 Calculate $\rho = x \cos \theta + y \sin \theta$

 Accumulate $H(\rho, \theta) = H(\rho, \theta) + 1$

Find (ρ, θ) where $H(\rho, \theta)$ is a local maximum and larger than a threshold.

The detected lines are given by $\rho = x \cos \theta + y \sin \theta$.

Line detection by Hough transform

Algorithm

Initialise the bins $H(\rho, \theta)$ to all zeros.

For each edge point (x, y)

For θ from 0 to π

Calculate $\rho = x \cos \theta + y \sin \theta$


Accumulate $H(\rho, \theta) = H(\rho, \theta) + 1$

Find (ρ, θ) where $H(\rho, \theta)$ is a **local maximum** and larger than a **threshold**.



The detected lines are given by $\rho = x \cos \theta + y \sin \theta$.

- Why local maximum?
 - Similar as non-maximum suppression in edge detection.
- Why thresholding?
 - A few random points would not lead to a line being detected.

Input image



Canny edge
detection



Lines detected by
Hough transform

Hough transform


- In model fitting, (m, b) are estimated by minimising the fitting error

$$\min_{m,b} \sum_i [y_i - (mx_i + b)]^2$$


- Only one line will be detected.
- On the contrary, Hough transform can simultaneously detect multiple lines, as long as they are local maxima above a threshold.

Hough transform


- It can detect multiple lines simultaneously.




Input image



Edge map




Parameter space




Detected lines

Hough transform


- It is robust to noise.
 - Edge map is often generated after image smoothing.
 - Broken edge points can still vote and contribute to line detection.




Input image



Edge map




Parameter space




Detected lines

Hough transform


- It is robust to object occlusion (e.g. the rectangle covered by a bunny).
 - The remaining edge points vote and contribute to line detection.




Input image



Edge map



Parameter space



Detected lines

Hough transform

- It is used not just for detecting lines, but also for other shapes, such as circles.
- We can parameterise a circle as,
$$(x - a)^2 + (y - b)^2 = r^2$$
 - The image space (x, y) is transformed to the parameter space (a, b, r) .
 - This is a very large search space (a lot of bins).
 - However, if we know the radius r , it would be easier.

Circle detection

$$(x - a)^2 + (y - b)^2 = r^2$$

- If the radius r is known, then for each edge point (x, y) , we only need to vote for possible values of (a, b) .
- It is still a circle in the parameter space $H(a, b)$.

$$(a - x)^2 + (b - y)^2 = r^2$$

Circle detection





Image space

$$(a - x)^2 + (b - y)^2 = 1, \text{ assume } r = 1$$



Parameter space

Vote result:
 $a = 2, b = 2$

Circle detection

$$(x - a)^2 + (y - b)^2 = r^2$$

- If the radius r is unknown, then it is a 3D parameter space $H(a, b, r)$.
- We set a range for the radius r .
For each $r \in [r_{min}, r_{max}]$
For each edge point (x, y)
We vote for possible values of (a, b) and accumulate $H(a, b, r)$.
- For example, we can start from $r = 1$ pixel to 10 pixels, each time increasing by 1 pixel.

Circle equations

- Standard form

$$(x - a)^2 + (y - b)^2 = r^2$$

- Parametric form using trigonometric functions

$$x = a + r \cdot \cos \theta$$

$$y = b + r \cdot \sin \theta$$

- This form gives us some ideas for acceleration.
- If we know the angle θ (direction) from the edge point (x, y) to the circle centre (a, b) , we can more accurately vote in the parameter space.

Circle detection





Image space

$$x = a + r \cdot \cos \theta$$
$$y = b + r \cdot \sin \theta$$



Parameter space

If we do not know θ , we vote to a whole circle.

Circle detection





Image space

$$x = a + r \cdot \cos \theta$$
$$y = b + r \cdot \sin \theta$$



Parameter space

If we know θ ,
we will only
vote along this
direction.

Circle detection





Image space

Since this is an edge point, we know its direction.

$$x = a + r \cdot \cos \theta$$
$$y = b + r \cdot \sin \theta$$



Parameter space

If we know θ , we will only vote along this direction.


Circle detection

- Parametric form using trigonometric functions

$$\begin{aligned}x &= a + r \cdot \cos \theta \\y &= b + r \cdot \sin \theta\end{aligned}$$

- The edge point (x, y) comes from an edge detection algorithm, so of course we know its direction θ .
- This narrows down our voting area in the parameter space $H(a, b, r)$. We simply move long θ (or opposite θ) for a distance of r .

Circle detection



We can assume that the edge direction θ is measured to accuracy of $\pm\phi$ and vote within this shaded area.

Circle detection by Hough transform

Algorithm

Initialise the bins $H(a, b, r)$ to all zeros.

For each possible radius $r \in [r_{min}, r_{max}]$

 For each edge point (x, y)

 Let θ to be gradient direction, or opposite gradient direction

 Calculate $a = x - r \cdot \cos \theta$, $b = y - r \cdot \sin \theta$

 Accumulate $H(a, b, r) = H(a, b, r) + 1$

Find (a, b, r) where $H(a, b, r)$ is a local maximum and larger than a threshold.

The detected circles are given by $x = a + r \cdot \cos \theta$, $y = b + r \cdot \sin \theta$.





Image of the moon surface, where Apollo 16 landed.



Circle detection by Hough transform

Hough transform

- Advantages
 - It detects multiple instances.
 - Robust to image noise.
 - Robust to occlusion.
- Limitations
 - The computational complexity can be high. For each edge point, we need to vote to a 2D or even 3D parameter space.
 - We need to carefully set some parameters, such as the parameters for the edge detector, the threshold for the accumulator or the range of circle radius.

Hough transform

- Apart from lines and circles, we can also use Hough transform for detecting other shapes.

- Ellipses

$$\frac{(x - c)^2}{a^2} + \frac{(y - d)^2}{b^2} = 1$$

- Planes in a 3D space

$$\frac{x}{a} + \frac{y}{b} + \frac{z}{c} = 1$$

- Other shapes that can be analytically represented

Hough transform

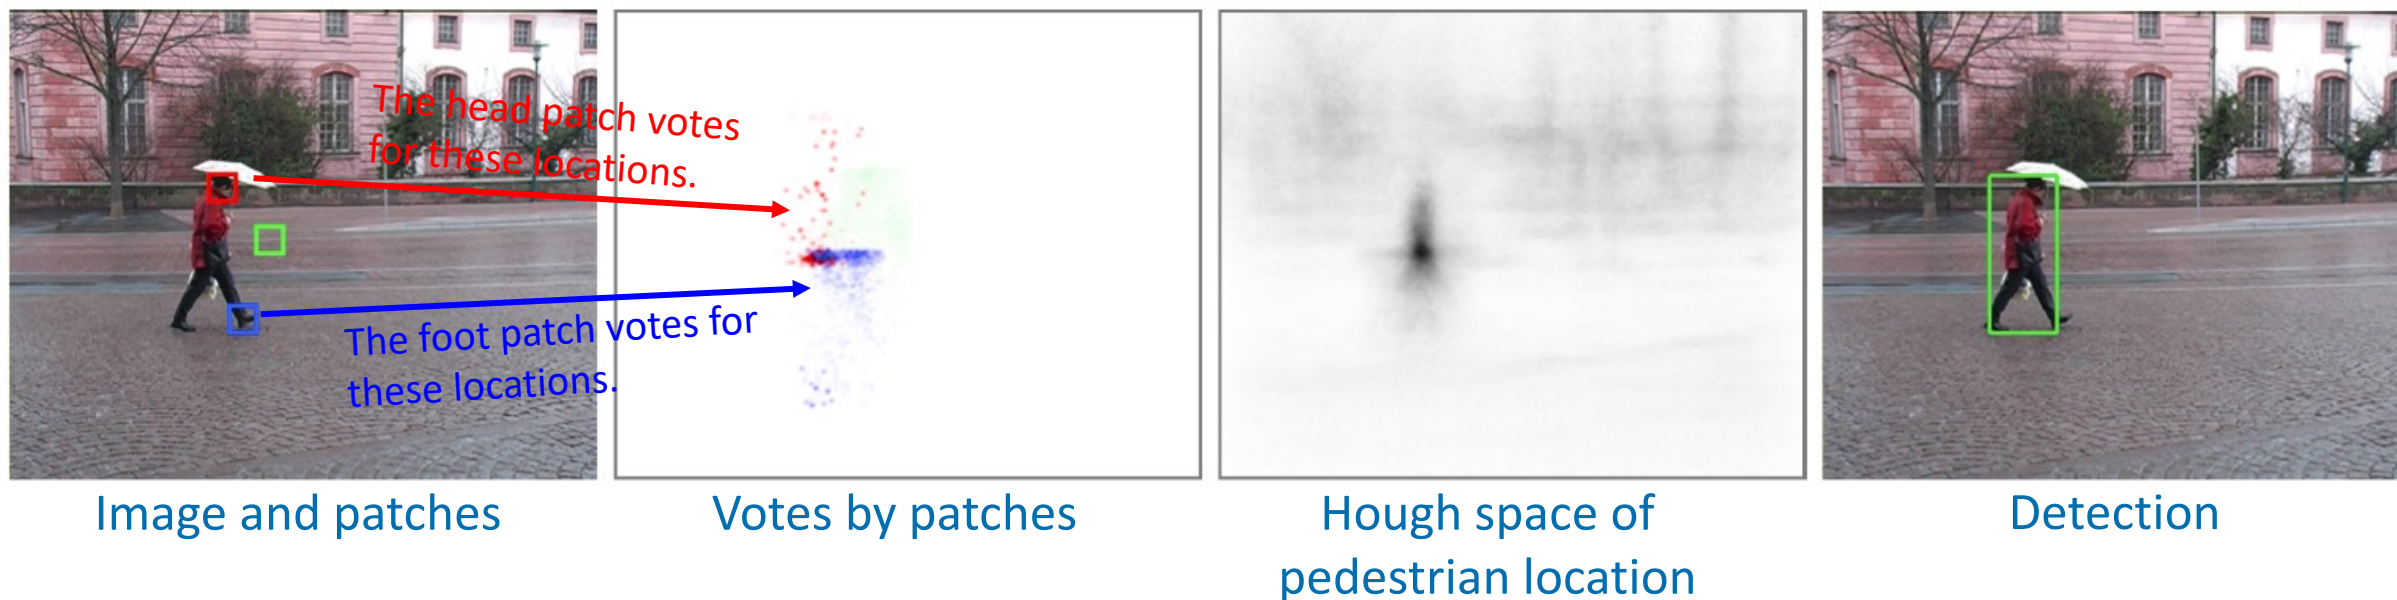
- When we vote in the parameter space, we can add some weights.
- Instead of using equal vote for each edge point, the vote can be weighted by the gradient magnitude so stronger edge points get higher weights.

Generalise the Hough transform idea

- In Hough transform, our aim is to detect some shapes or objects.
- There are two spaces, the image space and the parameter space (Hough space).
 - If the shape can be described by some parameters using an analytical equation, we use this equation to perform voting to the parameter space.
 - If it is not a simple shape without an analytical equation, as long as we have a model to describe it, we can still vote.

Hough forests for pedestrian detection

- The vote is performed by a machine learning model (random forest).
- The model predicts a displacement vector from the patch centre, given the image feature of the patch.



Hough transform

- Hough transform is an image analysis technique to detect shapes by a voting procedure.
 - Line detection
 - Circle detection
 - General shapes or patterns

References

- Sec. 4.3.2 Hough transforms. Richard Szeliski, Computer Vision: Algorithms and Applications (<http://szeliski.org/Book>).