

Motion II

Dr Wenjia Bai

Department of Computing & Brain Sciences

Motion

- Optic flow methods estimate motion for each pixel in the image.
- In this lecture, we will talk about
 - How to estimate motion for objects of interest?
 - How to recognise actions in videos?



Optic flow methods aim to estimate a dense (pixel-wise) motion field.



Object tracking aims to estimate the motion for one or multiple objects in a video.

Object tracking methods

- Lucas-Kanade tracker
- Correlation filter
- Tracking-by-detection

Lucas-Kanade tracker

- The Lucas-Kanade method is a general framework for both optic flow estimation and object tracking.
- Basic assumptions
 - Constant brightness
 - Small motion

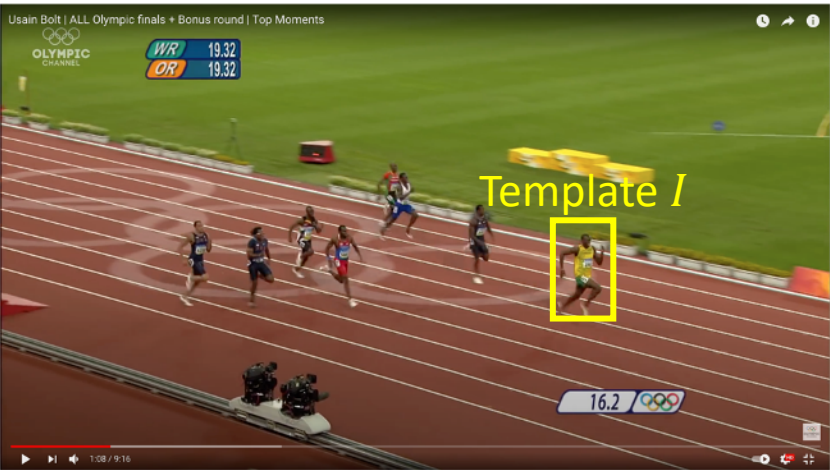
Lucas-Kanade tracker

- Lucas-Kanade aims to estimate the motion from the template image I to the image J in the next time frame.
- With the brightness constancy constraint, we formulate the following optimisation problem,


$$\min_{u,v} E(u, v) = \sum_x \sum_y [I(x, y) - J(x + u, y + v)]^2$$

x, y : pixels in the
template image

u, v : motion from template I to image J .
Point (x, y) on I corresponds to Point
 $(x + u, y + v)$ on J .



time t



time $t + 1$



time $t + 2$

Optimisation

- Cost function

$$\min_{u,v} E(u, v) = \sum_x \sum_y [I(x, y) - J(x + u, y + v)]^2$$

- Differentiate E with respect to u, v and let the derivatives be 0,

$$\frac{\partial E}{\partial u} = -2 \sum_x \sum_y [I(x, y) - J(x + u, y + v)] \frac{\partial J}{\partial x} = 0$$

$$\frac{\partial E}{\partial v} = -2 \sum_x \sum_y [I(x, y) - J(x + u, y + v)] \frac{\partial J}{\partial y} = 0$$

Optimisation

- With small motion assumption and Taylor expansion for J , we have

$$\frac{\partial E}{\partial u} = -2 \sum_x \sum_y \left[I(x, y) - J(x, y) - \frac{\partial J}{\partial x} u - \frac{\partial J}{\partial y} v \right] \frac{\partial J}{\partial x} = 0$$
$$\frac{\partial E}{\partial v} = -2 \sum_x \sum_y \left[I(x, y) - J(x, y) - \frac{\partial J}{\partial x} u - \frac{\partial J}{\partial y} v \right] \frac{\partial J}{\partial y} = 0$$

- With small motion assumption, we approximate $\frac{\partial J}{\partial x}$ by $\frac{\partial I}{\partial x}$, $\frac{\partial J}{\partial y}$ by $\frac{\partial I}{\partial y}$.
- We also have $\frac{\partial I}{\partial t} = J(x, y) - I(x, y)$.

Optimisation

- We can rewrite the equations as,

$$\frac{\partial E}{\partial u} = -2 \sum_x \sum_y \left[-\frac{\partial I}{\partial t} - \frac{\partial I}{\partial x} u - \frac{\partial I}{\partial y} v \right] \frac{\partial I}{\partial x} = 0$$
$$\frac{\partial E}{\partial v} = -2 \sum_x \sum_y \left[-\frac{\partial I}{\partial t} - \frac{\partial I}{\partial x} u - \frac{\partial I}{\partial y} v \right] \frac{\partial I}{\partial y} = 0$$

- Or simply as,

$$\frac{\partial E}{\partial u} = -2 \sum_x \sum_y [-I_t - I_x u - I_y v] I_x = 0$$
$$\frac{\partial E}{\partial v} = -2 \sum_x \sum_y [-I_t - I_x u - I_y v] I_y = 0$$

Optimisation

- Rewrite

$$\frac{\partial E}{\partial u} = -2 \sum_x \sum_y [-I_t - I_x u - I_y v] I_x = 0$$
$$\frac{\partial E}{\partial v} = -2 \sum_x \sum_y [-I_t - I_x u - I_y v] I_y = 0$$

in matrix form, we have

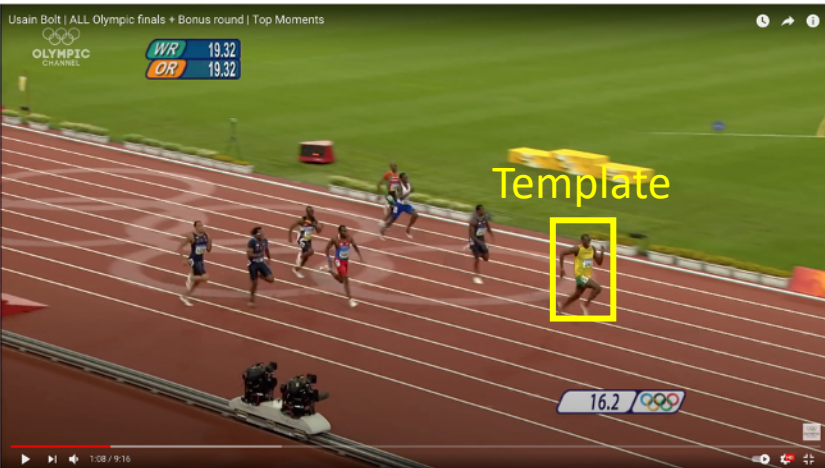
$$-\sum_x \sum_y \begin{bmatrix} I_x I_t \\ I_y I_t \end{bmatrix} - \sum_x \sum_y \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = 0$$

Optimisation

- Therefore the motion (u, v) can be solved,

$$\begin{bmatrix} u \\ v \end{bmatrix} = - \left(\sum_x \sum_y \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \right)^{-1} \sum_x \sum_y \begin{bmatrix} I_x I_t \\ I_y I_t \end{bmatrix}$$


- We get the same solution as the Lucas-Kanade optic flow method, if you check the slides from last lecture.
 - Lucas-Kanade optic flow calculates matrix by summing over a small neighbourhood.
 - Lucas-Kanade tracker calculates matrix by summing over pixels within the template image.



time t



time $t + 1$



time $t + 2$

Lucas-Kanade tracker

- In the previous derivation, we assume the motion is simply translation,

$$\min_{u,v} E(u, v) = \sum_x \sum_y [I(x, y) - J(x + u, y + v)]^2$$

- For general cases, we use a parametric model for the motion,

$$\min_{u,v} E(u, v) = \sum_x \sum_y [I(x, y) - J(W(x, y; p)))]^2$$

- For example,

$$W(x, y; p) = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

- The problem can be solved similarly.

Object tracking

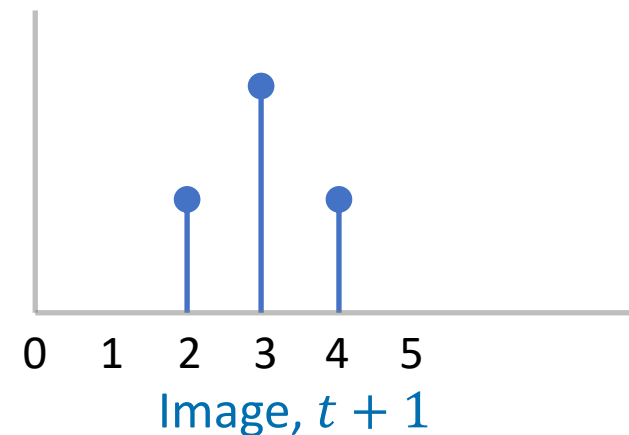

- The Lucas-Kanade method is a classical method.
 - However, the brightness constancy assumption may not always hold.
 - Lucas-Kanade only uses pixel intensities. It does not learn discriminative features for the template.
- Some recent object tracking methods may perform better.
 - Correlation filter method
 - Tracking by detection

Correlation filter

- We aim to maximise the correlation between template features and image features in the next time frame.

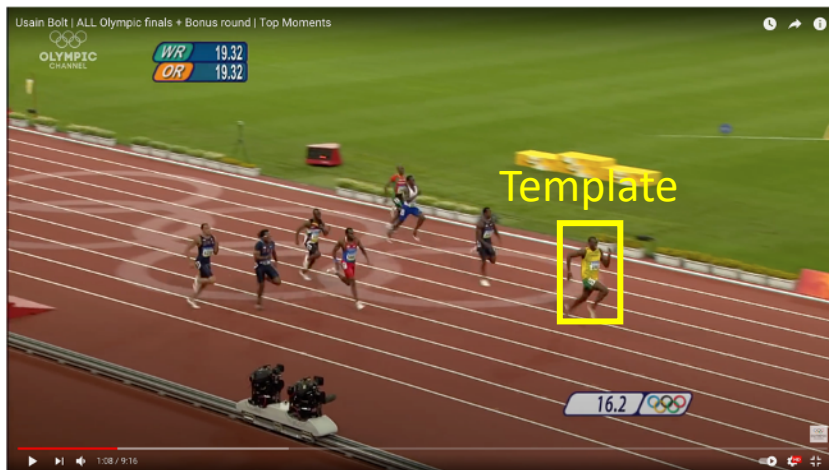
$$(f \star h)[n] = \sum_{m=-\infty}^{\infty} f[m]h[n + m]$$

- Correlation can be more robust to illumination changes than sum of squared difference.

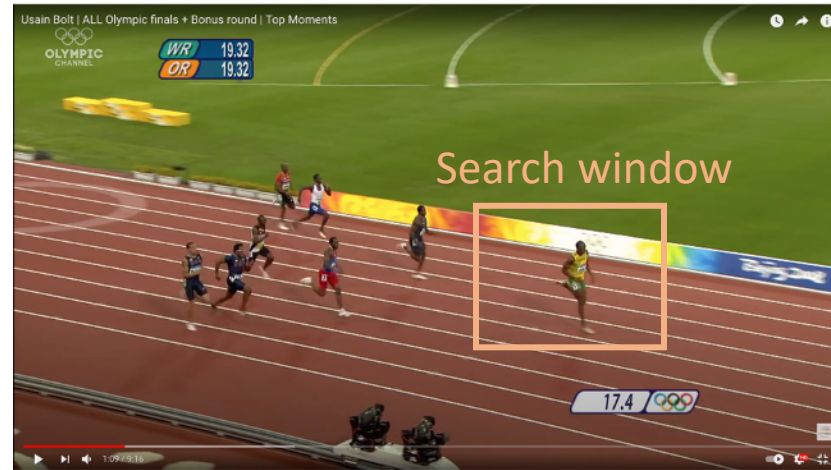


Correlation filter

- For 2D images, we can find where the correlation achieves the maximum between the template features and features in a search window.
- The features can be learnt from CNNs.



time t




time $t + 1$



time $t + 2$

Correlation filter

- A recent approach uses a Siamese network to learn and compare features.



AlexNet


[224x224x3] **Input**
[55x55x96] **Conv1**, 11x11, 96, s=4
[55x55x96] **Norm1**
[27x27x96] **Pool1**
[27x27x256] **Conv2**, 5x5, 256
[27x27x256] **Norm2**
[13x13x256] **Pool2**
[13x13x384] **Conv3**, 3x3, 384
[13x13x384] **Conv4**, 3x3, 384
[13x13x256] Conv5, 3x3, 256 convolutional
[13x13x256] **Norm3** feature map
[6x6x256] **Pool3**
[4096] **FC1**
[4096] **FC2**
[1000] **FC3** (class score)

VGG-16

[224x224x3] **Input**
[224x224x64] 3x3 conv1, 64
[224x224x64] 3x3 conv1, 64
[112x112x64] **Pool**
[112x112x128] 3x3 conv2, 128
[112x112x128] 3x3 conv2, 128
[56x56x128] **Pool**
[56x56x256] 3x3 conv3, 256
[56x56x256] 3x3 conv3, 256
[56x56x256] 3x3 conv3, 256
[28x28x256] **Pool**
[28x28x512] 3x3 conv4, 512
[28x28x512] 3x3 conv4, 512
[28x28x512] 3x3 conv4, 512
[14x14x512] **Pool**
[14x14x512] 3x3 conv5, 512
[14x14x512] 3x3 conv5, 512
[14x14x512] 3x3 conv5, 512 convolutional
[7x7x512] **Pool** feature map
[4096] **FC, 4096**
[4096] **FC, 4096**
[1000] **FC, 1000**


Correlation filter

- The network parameters are trained to predict a ground truth score map.




Training data

Template



Search window



The ground truth score is +1 within the yellow circle and -1 outside.



Example tracking results.

Tracking by detection

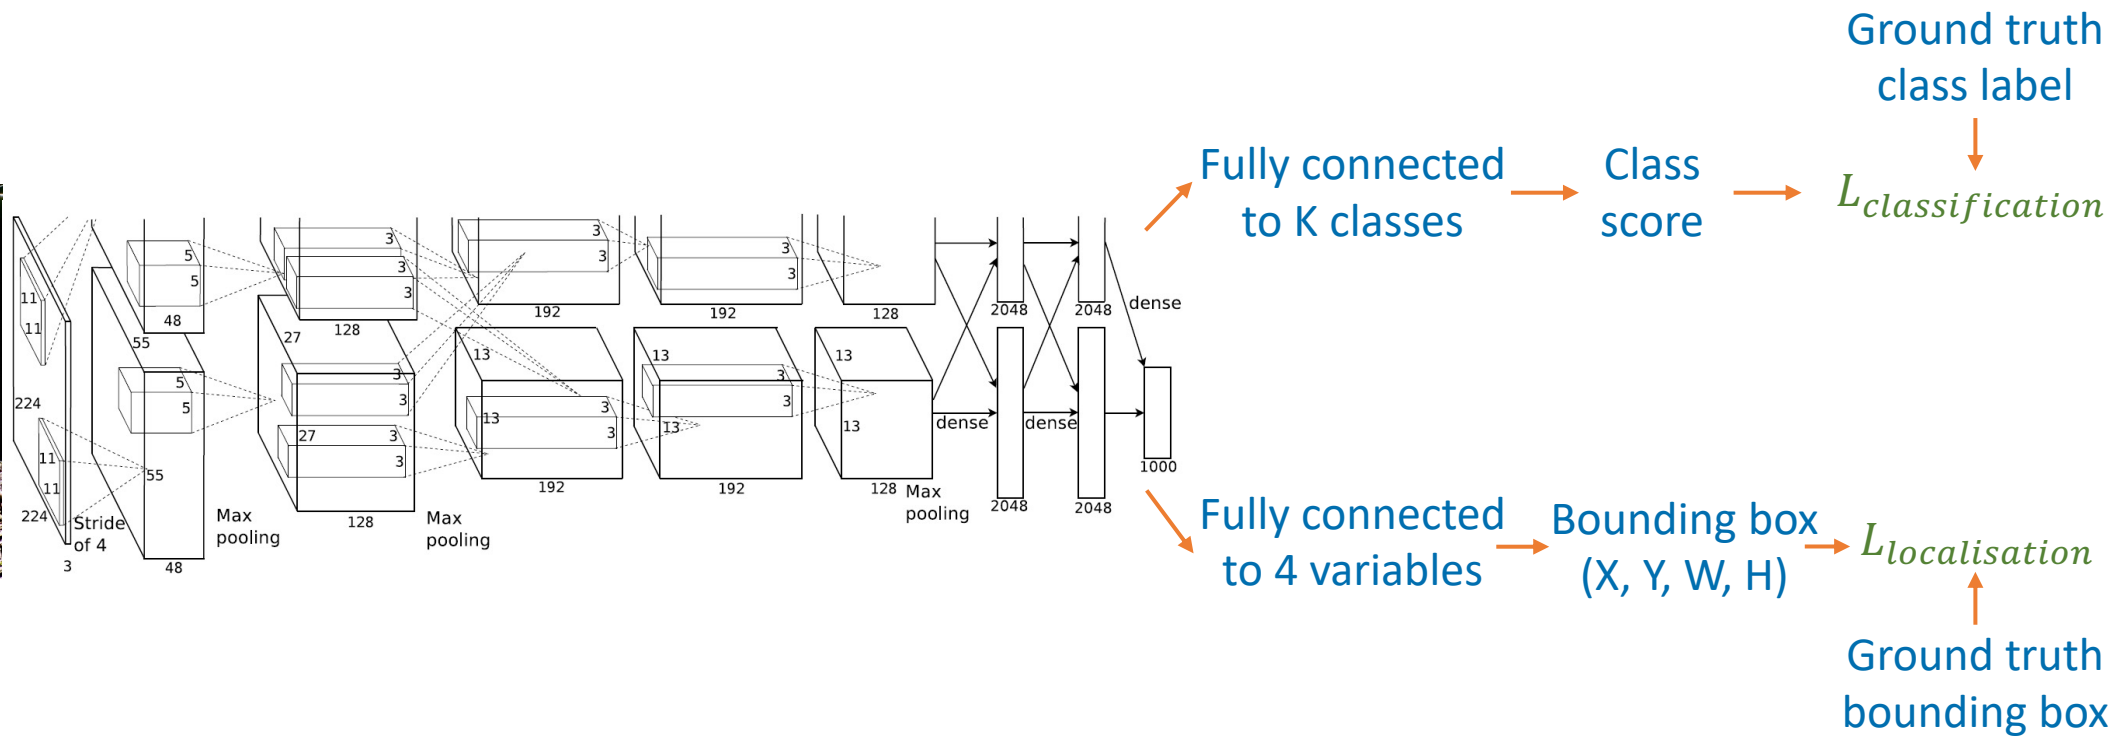
- Isn't object tracking a bit like object detection?
 - Yes, we can perform tracking by detection.
 - We want to learn what the object of interest looks like.
- Apart from their similarity, there are also some differences between object tracking and object detection.
 - Tracking concerns about videos.
 - In tracking, we have an initial bounding box. We are only interested in how this object moves. We do not need to detect every objects.

Tracking by detection

- Since we are interested in the object (initial bounding box) defined in the first time frame, we can learn features specific for this object.
- We can perform online learning.
 - Online: using information from this video, while it is streamed.
 - Offline: using information from an existing training set.

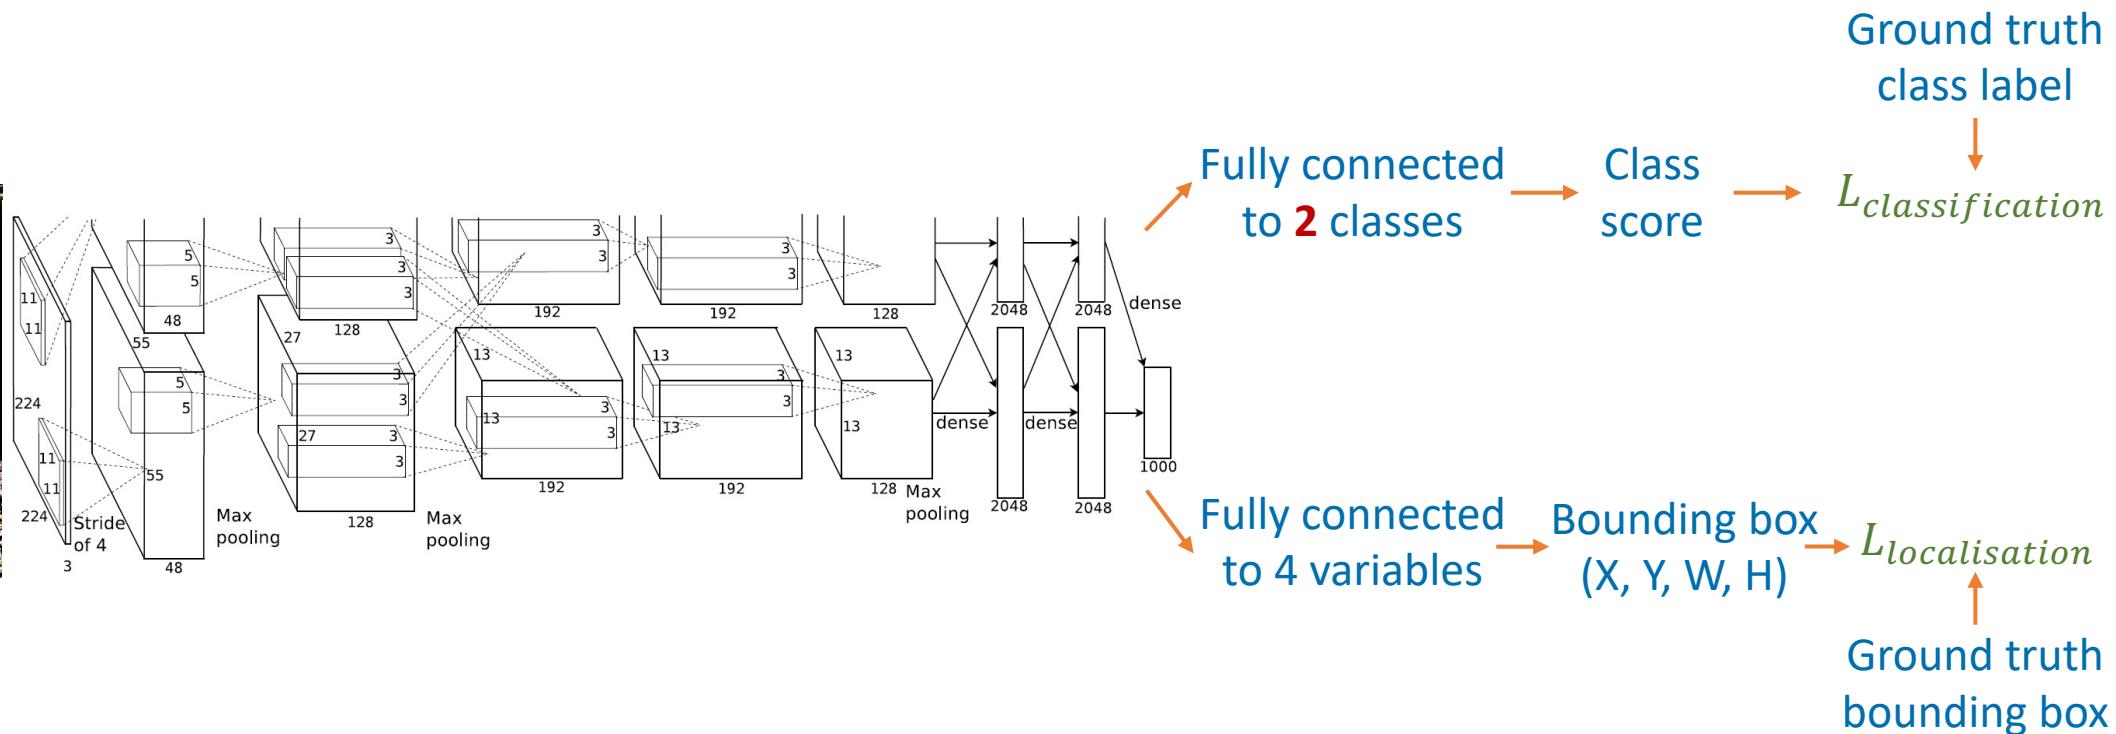
Object detection

- At each sliding window, we perform two tasks:
 - Task 1: classification
 - Task 2: localisation




Tracking by detection

- Using the same framework as object detection
 - However, we only perform binary classification, instead of multi-class classification.
 - The localisation task enables us to adjust the size of the object during tracking.



Training data

- The positive and negative samples for binary classification are extracted from the first time frame. More samples can be generated by
 - Data augmentation (translation, rotation, etc).
 - Using more time frames while the video is streamed.



time t



time $t + 1$




time $t + 2$

Training data


- Most layers of the convolutional network are pre-trained offline using large datasets such as ImageNet.
- Only the last few layers are trained online, using dozens or hundreds of samples extracted from the video.



Samples from
video



Pre-trained
feature extractor



Online trained
classifier and localiser

Tracking by detection



Tracking by detection



Action recognition

- Apart from object tracking, action recognition is another interesting application that involves videos and motion.



It is challenging to recognise actions from still images.

Action recognition

- How do we curate a dataset for action recognition?
- Kinetics Human Action Video Dataset [1,2]
 - 600 human action classes
 - More than 600 examples per class
 - Each from a unique YouTube video




It is challenging to recognise actions from still images.

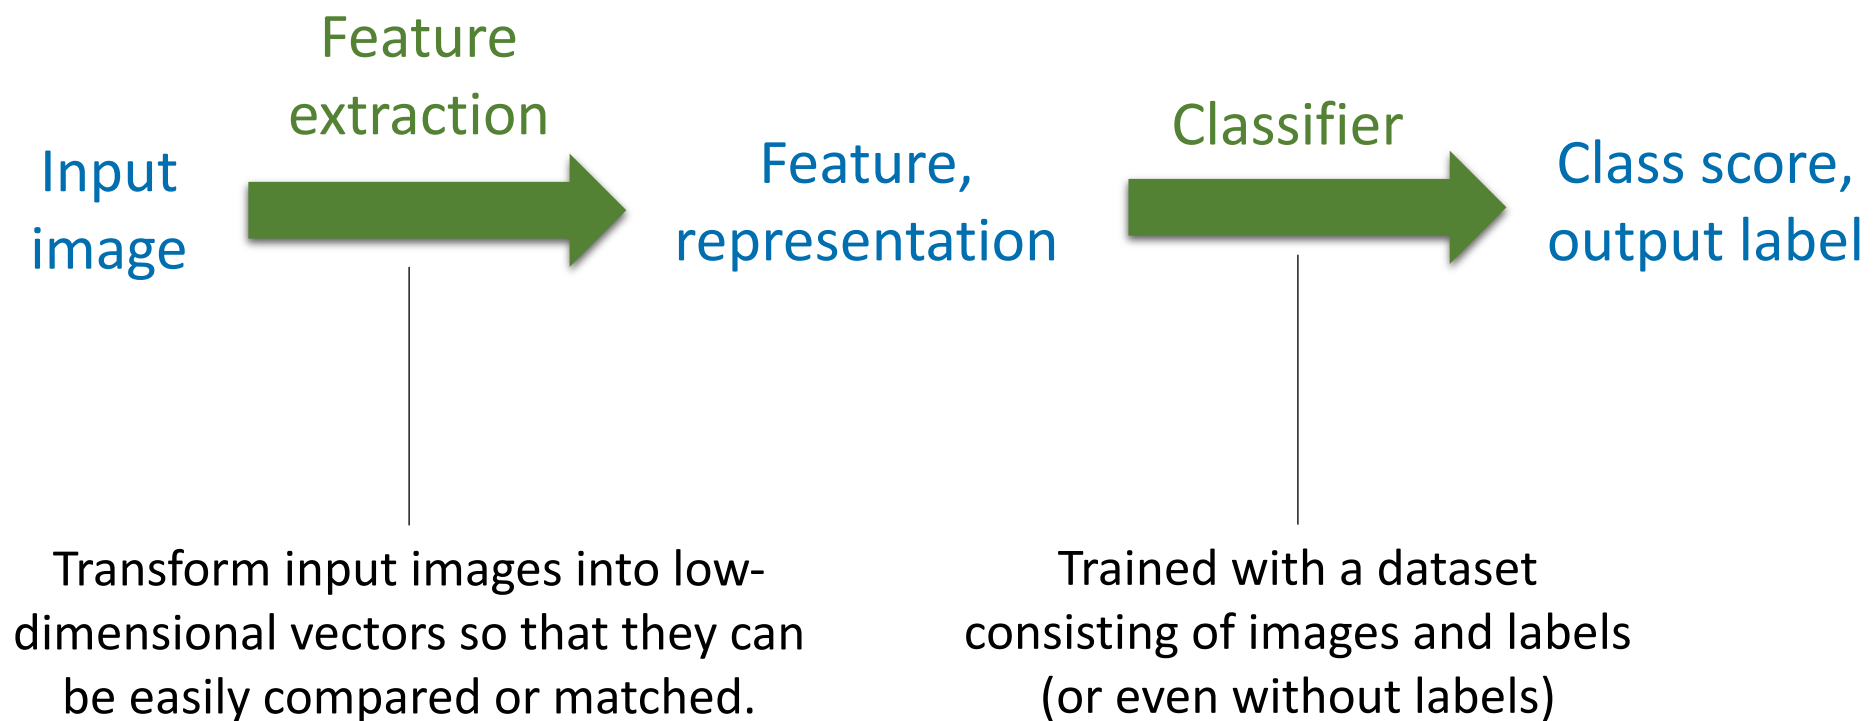
[1] J. Carreira and A. Zisserman. Quo Vadis, action recognition? A new model and the Kinetics dataset. CVPR 2017.

[2] <https://deepmind.com/research/open-source/open-source-datasets/kinetics/>

DEEPMIND
KINETICS VIEWER

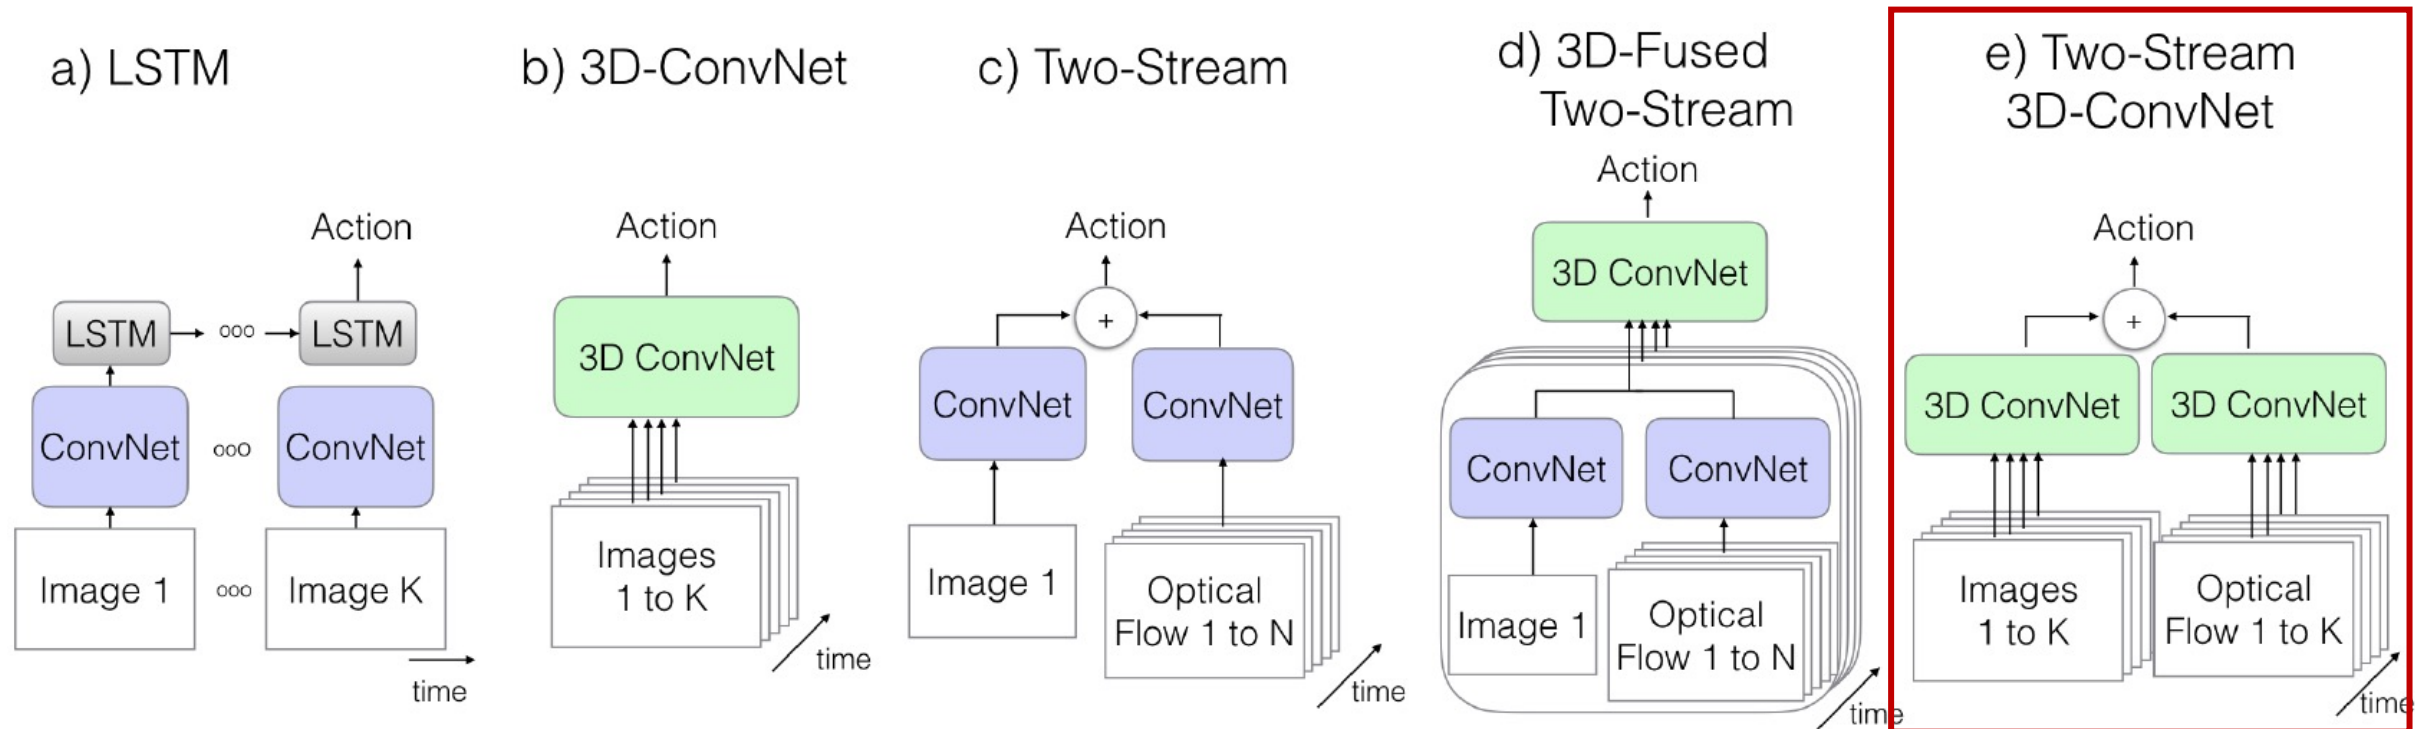


How does image classification work?



How do we recognise actions?

- We use both image and motion information.
 - Image: static features extracted from a single image
 - Motion: dynamic features extracted from videos, e.g. optic flow fields
- We can combine both to train a classifier for action recognition.



Network architectures for action recognition, which use both image and motion information.

Architecture	UCF-101			HMDB-51			miniKinetics		
	RGB	Flow	RGB + Flow	RGB	Flow	RGB + Flow	RGB	Flow	RGB + Flow
(a) LSTM	81.0	–	–	36.0	–	–	69.9	–	–
(b) 3D-ConvNet	51.6	–	–	24.3	–	–	60.0	–	–
(c) Two-Stream	83.6	85.6	91.2	43.2	56.3	58.3	70.1	58.4	72.9
(d) 3D-Fused	83.2	85.8	89.3	49.2	55.5	56.8	71.4	61.0	74.0
(e) Two-Stream I3D	84.5	90.6	93.4	49.8	61.9	66.4	74.1	69.6	78.7

The action recognition accuracy is improved by using both image and motion information, compared to using image information alone.

Summary

- Object tracking
 - Lucas-Kanade tracker
 - Correlation filter
 - Tracking-by-detection
- Action recognition